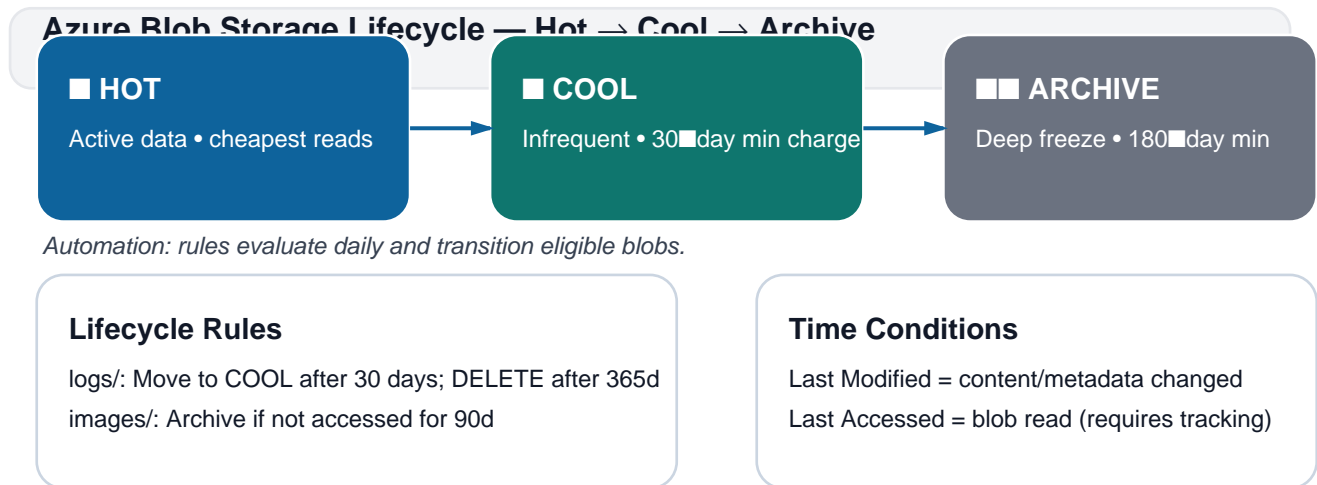


Azure Storage: Hot vs Cool vs Archive + Lifecycle Rules

■ **Built an Azure Storage lab** that auto-moves data between tiers and trims costs. Here's the 60-second story you can paste on LinkedIn—plus a clean diagram you can screenshot.



■ Plain-English tiers

■ **Hot** — for files you open a lot (cheapest to read, pricier to store)

■ **Cool** — for infrequent access (cheaper to store, 30-day minimum charge)

■ **Archive** — deep freeze for long retention (180-day minimum; rehydrate to read)

■ What I built

- Storage Account with secure defaults (private containers, soft delete, optional versioning)
- Two containers: logs/ and images/
- Lifecycle rules to: move logs/ → COOL after 30d, delete after 365d; archive images/ after 90d of no access

■ Why it matters

- Cuts storage costs automatically
- Mirrors real usage (hot vs cold data)
- Reinforces AZ900 fundamentals (governance, cost control, lifecycle)

■ Gotchas

- Cool has a 30-day minimum; Archive has 180-day
- Archive blobs are offline until rehydrated (time + cost)

➔ Next up

- Turn on access-time tracking across more containers
- Add alerts to flag unexpected storage growth

■ Paste-ready caption

■ I just built an Azure Storage lab (Hot vs Cool vs Archive) with auto lifecycle rules—here's what I learned ■

Most teams pay for storage they don't actually use. I set up **Azure Blob Storage** with **lifecycle policies** that move old data to cheaper tiers—and eventually delete it—*automatically*.

Hot = active files (cheapest reads) • **Cool** = infrequent (30■day min) • **Archive** = deep freeze (180■day min; rehydrate to read).

Built: secure Storage Account, containers (logs/, images/), and rules: logs→Cool after 30d (delete 365d), images→Archive after 90d of no access.

Why: ■ cost control • ■ real-world lifecycle • ■ AZ■900 fundamentals.

Gotchas: Cool/Archive minimums; Archive is offline until rehydrated.

#Azure #AZ900 #Cloud #Storage #CostOptimization #DevOps #LearningInPublic